

Machine learning foundations I

Mitko Veta

Eindhoven University of Technology
Department of Biomedical Engineering

2019

Learning goals

At the end of this lecture you will:

- ▶ Have an understanding of the goal of machine learning (ML) models.
- ▶ Have a good understanding of basic mathematical concepts used in ML and be able to apply them in the design and implementation of ML methods.

Overview

Topics covered in this lecture:

1. Linear algebra
2. Gradient-based optimization
3. Two simple machine learning models
 - Linear model
 - Nearest-neighbours model
4. Probability theory

Linear algebra

Materials:

- ▶ Chapter 1.2 from **deeplearning**
- ▶ **linearalgebra**

Matrices

- ▶ A matrix is a 2-D array of numbers

$$\begin{bmatrix} a_{1,1} & a_{1,2} \\ a_{2,1} & a_{2,2} \end{bmatrix}$$

- ▶ Example notation for type and shape
 $\mathbf{A} \in \mathbb{R}^{m \times n}$

Tensors

- ▶ A tensor is an array of numbers that may have
 - ▶ a zero dimensions and be a scalar,
 - ▶ one dimension and be a vector,
 - ▶ two dimensions and be a matrix,
 - ▶ more dimensions ...

Side note: One of the most popular frameworks for implementing deep machine learning models is called TensorFlow (<https://www.tensorflow.org/>).

Transpose matrix

$$(\mathbf{A}^T)_{i,j} = \mathbf{A}_{j,i}$$

$$\mathbf{A} = \begin{bmatrix} A_{1,1} & A_{1,2} \\ A_{2,1} & A_{2,2} \\ A_{3,1} & A_{3,2} \end{bmatrix} \Rightarrow \begin{bmatrix} A_{1,1} & A_{2,1} & A_{3,1} \\ A_{1,2} & A_{2,2} & A_{3,2} \end{bmatrix}$$

The transpose matrix is a mirror image with regard to the main diagonal

$$(\mathbf{AB})^T = \mathbf{B}^T \mathbf{A}^T$$

Matrix (dot) product

$$\mathbf{C} = \mathbf{AB}$$

The matrices must be compatible: an $m \times n$ matrix is multiplied with an $n \times r$ matrix and as a result an $m \times r$ matrix is obtained

$$C_{i,j} = \sum_k A_{i,k} B_{k,j}$$

$$\mathbf{A} \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \\ 7 & 8 & 9 \\ 10 & 11 & 12 \end{bmatrix} \times \mathbf{B} \begin{bmatrix} 1 & 2 & 3 & 4 & 5 \\ 6 & 7 & 8 & 9 & 10 \\ 11 & 12 & 13 & 14 & 15 \end{bmatrix} = \mathbf{C} [4 \times 5]$$

$$C_{2,5} = A_{2,1}B_{1,5} + A_{2,2}B_{2,5} + A_{2,3}B_{3,5} = 4 \cdot 5 + 5 \cdot 10 + 6 \cdot 15 = 160$$

Matrix (dot) product

- ▶ In general matrix multiplication is not commutative, i.e., most of the time $\mathbf{AB} \neq \mathbf{BA}$.
- ▶ Depending on the dimensions sometimes \mathbf{AB} or \mathbf{BA} are not possible.
- ▶ As a special case the matrix can be a (column or row) vector; an $m \times n$ matrix is multiplied with a $n \times 1$ vector to obtain a $m \times 1$ vector.

Systems of equations

► $\mathbf{Ax} = \mathbf{b}$ expands to

$$\mathbf{A}_{1,:}\mathbf{x}_1 = \mathbf{b}_1 \quad (1)$$

$$\mathbf{A}_{2,:}\mathbf{x}_2 = \mathbf{b}_2 \quad (2)$$

$$\dots \quad (3)$$

$$\mathbf{A}_{m,:}\mathbf{x}_m = \mathbf{b}_m \quad (4)$$

Solving systems of linear equations

- ▶ A linear system of equations can have
 - ▶ no solutions,
 - ▶ many solutions,
 - ▶ exactly one solution.
- ▶ Only one solution implies that multiplication by a matrix is an invertible operation.

Matrix inversion

- ▶ Matrix inverse is defined with

$$\mathbf{A}^{-1}\mathbf{A} = \mathbf{I}_n$$

- ▶ A system of linear equations can be solved using inverse matrix

$$\mathbf{Ax} = \mathbf{b}$$

$$\mathbf{A}^{-1}\mathbf{Ax} = \mathbf{A}^{-1}\mathbf{b}$$

$$\mathbf{I}_n\mathbf{x} = \mathbf{A}^{-1}\mathbf{b}$$

$$\mathbf{x} = \mathbf{A}^{-1}\mathbf{b}$$

- ▶ This is useful mostly for abstract analysis.
- ▶ From a numerical point of view there are much more efficient methods.

Invertibility

A matrix cannot be inverted if

- ▶ the number of rows and columns is not the same, or
- ▶ some rows and columns are "redundant" ("linearly dependent", "low rank").

Norms

- ▶ Norms are functions that measure how "large" a vector is.
- ▶ Similar to a distance between zero and the point represented by the vector
 - ▶ $f(\mathbf{x}) = 0 \Rightarrow \mathbf{x} = 0$
 - ▶ $f(\mathbf{x} + \mathbf{y}) \leq f(\mathbf{x}) + f(\mathbf{y})$ (**the triangle inequality**)
 - ▶ $\forall \alpha \in \mathbb{R} : f(\alpha \mathbf{x}) = \alpha f(\mathbf{x})$

Norms

- ▶ L^p - norm

$$\|\mathbf{x}\|_p = \left(\sum_i |x_i|^p \right)^{\frac{1}{p}}$$

- ▶ Most popular L^2 -norm (for $p = 2$)
- ▶ L_1 -norm (for $p = 1$): $\|\mathbf{x}\|_1 = \sum_i |x_i|$
- ▶ Max norm (for infinite p): $\|\mathbf{x}\|_\infty = \max_i |x_i|$

Special vectors and matrices

- ▶ Unit vector $\|\mathbf{x}\|_n = 1$
- ▶ Symmetric matrix $\mathbf{A} = \mathbf{A}^T$
- ▶ Orthogonal matrix

$$\mathbf{A}\mathbf{A}^T = \mathbf{I} = \mathbf{A}^T\mathbf{A}$$

- ▶ It follows that for orthogonal matrices $\mathbf{A}^T = \mathbf{A}^{-1}$

Eigendecomposition

- ▶ Eigenvector and eigenvalue

$$\mathbf{A}\mathbf{x} = \lambda\mathbf{x}$$

- ▶ Eigendecomposition of a matrix

$$\mathbf{A} = \mathbf{V}\text{diag}(\lambda)\mathbf{V}^{-1}$$

where $\text{diag}(\lambda)$ is a diagonal matrix having the (scalar) eigenvalues λ as diagonal elements.

Eigendecomposition

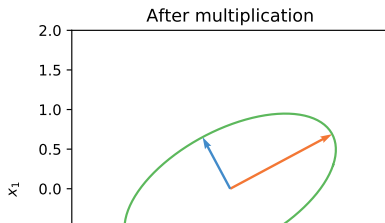
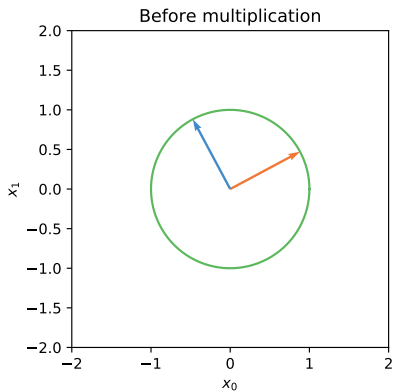
- ▶ Every real symmetric matrix has a real orthogonal eigendecomposition

$$\mathbf{A} = \mathbf{Q}\mathbf{\Lambda}\mathbf{Q}^T$$

where \mathbf{Q} is an orthogonal matrix composed of eigenvectors of \mathbf{A} and $\mathbf{\Lambda}$ is a diagonal matrix.

- ▶ The eigenvalue Λ_{ii} is associated with the eigenvector in column i of \mathbf{Q} , denoted as $\mathbf{Q}_{:,i}$.
- ▶ We can think of \mathbf{A} as scaling space by factor λ_i in the direction of its corresponding eigenvector $\mathbf{v}^{(i)}$ (represented by $\mathbf{Q}_{:,i}$).

Effect of eigenvalues



Singular value decomposition

- ▶ Similar to eigenvalue decomposition
- ▶ More general: matrix need not be square

$$\mathbf{A} = \mathbf{U}\mathbf{D}\mathbf{V}^T$$

- ▶ \mathbf{U} and \mathbf{V} are square matrices and are both orthogonal, \mathbf{D} is diagonal.
- ▶ The diagonal elements of \mathbf{D} are called **singular values** of matrix \mathbf{A} ; the columns of \mathbf{U} and \mathbf{V} are **left-singular** and **right-singular vectors** of \mathbf{A} , respectively.

Moore-Penrose pseudoinverse

- ▶ Matrix inversion is not defined on matrices that are not square.
- ▶ The **Moore-Penrose pseudoinverse** is defined as

$$\mathbf{A}^+ = \lim_{\alpha \searrow 0} (\mathbf{A}^T \mathbf{A} + \alpha \mathbf{I})^{-1} \mathbf{A}^T$$

Moore-Penrose pseudoinverse

Now we can consider

$$\mathbf{x} = \mathbf{A}^+ \mathbf{y}$$

- ▶ If the equation has
 - ▶ exactly one solution: this is the same as inverse,
 - ▶ no solution: gives the solution with the smallest error, $\|\mathbf{Ax} - \mathbf{y}\|_2$
 - ▶ many solutions: gives the solution with the smallest norm of \mathbf{x} .

Computing the pseudoinverse

- ▶ Efficient implementations are based on the formula allowed by the singular decomposition

$$\mathbf{A}^+ = \mathbf{V}\mathbf{D}^+ \mathbf{U}^T$$

- ▶ \mathbf{U} , \mathbf{D} , \mathbf{V} are from the singular value decomposition of \mathbf{A} .
- ▶ The pseudoinverse \mathbf{D}^+ of \mathbf{D} is obtained by taking the reciprocal non-zero elements and after that taking the transpose of the resulting matrix.

Gradient-based optimization

Materials:

- ▶ Chapters 1.4 and 1.5 from **deeplearning**
- ▶ **linearalgebra**

Gradient

- ▶ Let $f : \mathbb{R}^{m \times n} \mapsto \mathbb{R}$ be a function that takes $m \times n$ matrix \mathbf{A} as input and returns a real number (scalar).
- ▶ A **gradient** of f with respect to A is the matrix

$$\nabla_{\mathbf{A}} f(\mathbf{A}) = \begin{bmatrix} \frac{\partial f}{\partial A_{11}} & \frac{\partial f}{\partial A_{12}} & \cdots & \frac{\partial f}{\partial A_{1n}} \\ \frac{\partial f}{\partial A_{21}} & \frac{\partial f}{\partial A_{22}} & \cdots & \frac{\partial f}{\partial A_{2n}} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial f}{\partial A_{m1}} & \frac{\partial f}{\partial A_{m2}} & \cdots & \frac{\partial f}{\partial A_{mn}} \end{bmatrix}$$

- ▶ i.e. an $m \times n$ matrix with

$$(\nabla_{\mathbf{A}} f(\mathbf{A}))_{ij} = \frac{\partial f}{\partial A_{ij}}$$

- ▶ The size of the gradient of \mathbf{A} is the same as the size of A .

Gradient

- ▶ In the special case when A is a vector we obtain the (possibly more familiar) gradient

$$\nabla_{\mathbf{x}} f(\mathbf{x}) = \begin{bmatrix} \frac{\partial f}{\partial x_1} \\ \frac{\partial f}{\partial x_2} \\ \vdots \\ \frac{\partial f}{\partial x_m} \end{bmatrix}$$

- ▶ In general to define a gradient we require that the function returns a **real** value.

Jacobian

- ▶ The Jacobian \mathbf{J}_f is a generalization of the gradient for vector valued functions.
- ▶ Let $\mathbf{f} : \mathbb{R}^n \mapsto \mathbb{R}^m$ be a function that takes n -dimensional vector \mathbf{x} as input and returns a m -dimensional vector as an output.
- ▶ The Jacobian \mathbf{J}_f is defined as

$$\mathbf{J}_f = \begin{bmatrix} \frac{\partial f_1}{\partial x_1} & \frac{\partial f_1}{\partial x_2} & \cdots & \frac{\partial f_1}{\partial x_n} \\ \frac{\partial f_2}{\partial x_1} & \frac{\partial f_2}{\partial x_2} & \cdots & \frac{\partial f_2}{\partial x_n} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial f_m}{\partial x_1} & \frac{\partial f_m}{\partial x_2} & \cdots & \frac{\partial f_m}{\partial x_n} \end{bmatrix}$$

- ▶ Note that for the special case of a scalar-valued function, the Jacobian is the transpose of the gradient.

Optimization

- ▶ Most machine learning methods involve some kind of optimization.
 - ▶ One exception is the k -Nearest neighbour classifier introduced later.
- ▶ Optimization means minimizing or maximizing some function $f(\mathbf{x})$, i.e. finding the values of \mathbf{x} for which $f(\mathbf{x})$ has a minimum or a maximum.
- ▶ Notation: $\mathbf{x}^* = \operatorname{argmin} f(\mathbf{x})$

Gradient-based optimization

- ▶ The derivative tells us how to change x in order to make a small improvement of $f(x)$.
- ▶ Therefore, derivatives can be useful in optimization.

Two simple machine learning models

Materials:

- ▶ Chapter 2.3 from **elements**

The learning task

- ▶ Given a value of the input vector \mathbf{x} make a good prediction of the output y , denoted as \hat{y} .
- ▶ Both y and \hat{y} should take values from the same numerical set.
- ▶ Similarly, g and \hat{g} should both take values from the same set \mathcal{G} .
- ▶ We suppose that we have available a set of measurements (\mathbf{x}_i, y_i) or (\mathbf{x}_i, g_i) ($1 \leq i \leq N$) called **training data** (in matrix form: (\mathbf{X}, \mathbf{y}) and/or (\mathbf{X}, \mathbf{g})).
- ▶ Our task is to construct a prediction rule based on the training data.

The learning task

Example:

- ▶ **Variable values:** Let g (and therefore also \hat{g}) be two valued (categorical), e.g. $\mathcal{G} = \{\text{BLUE}, \text{ORANGE}\}$.
- ▶ **Encoding of g s with y s:** Then each class can be encoded binary, i.e., with $y \in \{0, 1\}$, e.g., **BLUE** and **ORANGE**, would correspond to 0 and 1, respectively.
- ▶ **Predicted output values:** \hat{y} ranges over the interval $[-\infty, +\infty]$ (of which $\{0, 1\}$ is a subset).
- ▶ **Prediction rule:** \hat{g} is assigned a (class label) **BLUE** if $\hat{y} < 0.5$ and **ORANGE**, otherwise.

Two simple approaches to prediction

- ▶ Linear model fit
 - ▶ strong assumptions about the structure of the decision boundary
- ▶ k -nearest neighbours
 - ▶ weak assumptions about the structure of the decision boundary

Linear model fit by least squares

- ▶ Despite relative simplicity one of the most important statistical tools
- ▶ Input vector $\mathbf{x}^T = (x_1, x_2, \dots, x_p)$
- ▶ Output y predicted using the model

$$\hat{y} = \hat{w}_0 + \sum_{j=1}^p x_j \hat{w}_j$$

- ▶ \hat{w}_i ($0 \leq i \leq p$) are the parameters of the linear model
- ▶ In vector form

$$\hat{y} = \hat{\mathbf{w}}^T \mathbf{x} = \mathbf{x}^T \hat{\mathbf{w}}$$

using the fact that the scalar (inner) product of two vectors is a commutative operation.

Linear model fit by least squares

- ▶ We assume that w_0 is in \mathbf{w} and 1 is included in \mathbf{x} .
- ▶ \hat{y} is a scalar, but in general can be a k -vector $\hat{\mathbf{y}}$, in which case \mathbf{w} becomes a $p \times k$ matrix of coefficients.

Linear model fit by least squares

- ▶ There are many ways to fit a linear model to a training dataset.
- ▶ **Least squares** method
 - ▶ We need to find coefficients \hat{w}_i which minimize the error estimated with the **residual sum of squares**

$$\text{RSS}(\mathbf{w}) = \sum_{i=1}^N (y_i - \mathbf{x}_i^T \mathbf{w})^2$$

assuming N input-output pairs.

- ▶ $\text{RSS}(\mathbf{w})$ is a quadratic function.
- ▶ A minimum always exists though not necessarily a unique one.

Linear model fit by least squares

- ▶ We look for the solution $\hat{\mathbf{w}}$ using the matrix notation:
- ▶ $\mathbf{y} = [y_1, y_2, \dots, y_N]^T$ is the vector formed from the N output vectors and \mathbf{X} is an $N \times p$ matrix

$$\text{RSS}(\mathbf{w}) = (\mathbf{y} - \mathbf{X}\mathbf{w})^T(\mathbf{y} - \mathbf{X}\mathbf{w})$$

- ▶ To find the minimum we differentiate with respect to \mathbf{w} which gives

$$(-\mathbf{X})^T(\mathbf{y} - \mathbf{X}\mathbf{w}) + (\mathbf{y} - \mathbf{X}\mathbf{w})^T(-\mathbf{X})$$

using the rule $(\mathbf{AB})^T = \mathbf{B}^T \mathbf{A}^T$ this is equivalent to

$$-2\mathbf{X}^T(\mathbf{y} - \mathbf{X}\mathbf{w})$$

Linear model fit by least squares

- ▶ To find the minimum our derivative must be $\mathbf{0}$, hence:

$$\mathbf{X}^T(\mathbf{y} - \mathbf{X}\mathbf{w}) = \mathbf{0}$$

$$\mathbf{X}^T\mathbf{y} - \mathbf{X}^T\mathbf{X}\mathbf{w} = \mathbf{0}$$

$$\mathbf{X}^T\mathbf{y} = \mathbf{X}^T\mathbf{X}\mathbf{w}$$

- ▶ If $\mathbf{X}^T\mathbf{X}$ is non-singular there exists a unique solution given by

$$\hat{\mathbf{w}} = (\mathbf{X}^T\mathbf{X})^{-1}\mathbf{X}^T\mathbf{y}$$

Question: Why not simply $\mathbf{y} - \mathbf{X}\mathbf{w} = \mathbf{0} \rightarrow \mathbf{y} = \mathbf{X}\mathbf{w} \rightarrow \hat{\mathbf{w}} = \mathbf{X}^{-1}\mathbf{y}$?

Linear model fit by least squares

- ▶ For each input \mathbf{x}_i there corresponds the fitted output

$$\hat{y}_i = \hat{y}_i(\mathbf{x}_i) = \hat{\mathbf{w}}^T \mathbf{x}_i$$

.

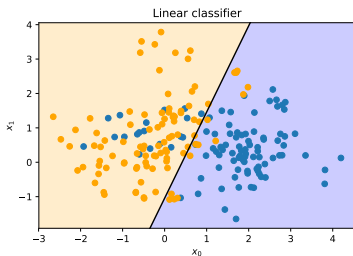
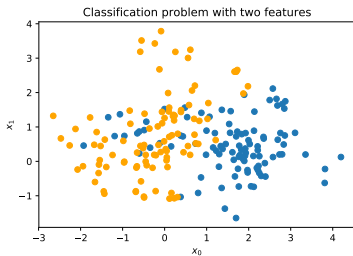
- ▶ This is called “making a prediction” for \mathbf{x}_i .
- ▶ The entire fitted surface (hyperplane) is fully characterized by the parameter vector $\hat{\mathbf{w}}$.
- ▶ After fitting the model, we can “discard” the training dataset.

Example: Linear model fit by least squares

- ▶ Scatter plot of training data on a pair of inputs x_1 and x_2
- ▶ Output class variable g has two values **BLUE** and **ORANGE**.
- ▶ Linear regression model fitted with the response variable y coded as 0 for **BLUE** and 1 for **ORANGE**.
- ▶ Fitted values \hat{y} converted to a fitted class variable \hat{g} as

$$\hat{g} = \begin{cases} \text{BLUE} & \text{if } \hat{y} \leq 0.5 \\ \text{ORANGE} & \text{if } \hat{y} > 0.5 \end{cases}$$

Example: Linear model fit by least squares



Example: Linear model fit by least squares

- ▶ Wrong classifications on both sides of the boundary
- ▶ Are the errors caused by the model or are they unavoidable?
- ▶ Two possible scenarios
 - ▶ **Scenario 1:** data generated from bivariate Gaussian distribution
 - ▶ **Scenario 2:** data generated from 10 Gaussian distributions; the means of these distributions are also distributed as Gaussian
- ▶ In Scenario 1 the linear boundary is the best we can do since the overlap is inevitable.
- ▶ In Scenario 2 the linear boundary is unlikely to be optimal (in fact the boundary is non-linear and disjoint).

Nearest-neighbours model

- ▶ In nearest-neighbour methods $\hat{y}(\mathbf{x})$ is determined based on the inputs (points) in the training set \mathcal{T} which are "closest" to the input \mathbf{x} .
- ▶ k -nearest neighbour fit is defined as

$$\hat{y}(\mathbf{x}) = \frac{1}{k} \sum_{\mathbf{x}_i \in N_k(\mathbf{x})} y_i$$

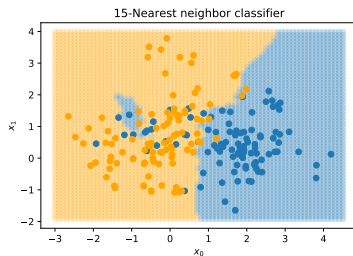
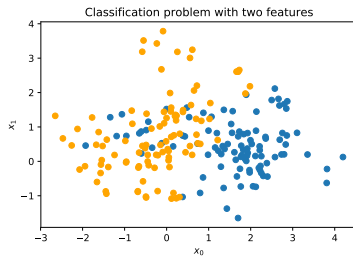
where $N_k(\mathbf{x})$ is the neighbourhood of \mathbf{x} consisting of the k "closest" points to \mathbf{x} .

- ▶ "Closeness" requires a definition of **metrics**.
- ▶ For the moment we assume Euclidian distance (each \mathbf{x} is a point in the hyperspace).
- ▶ An average of the classes of the k closest points

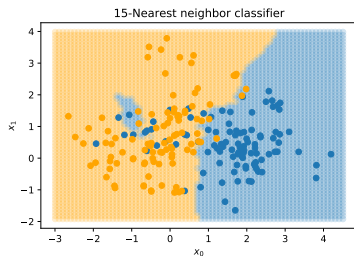
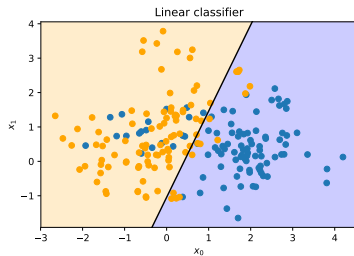
Back to the BLUE and ORANGE example

- ▶ We use the same training data as in the linear model example.
- ▶ New borderline between the classes generated with 15-nearest-neighbour model
- ▶ Since ORANGE is encoded as 1 \hat{y} is the proportion of ORANGE points in the 15-neighbourhood
- ▶ Class ORANGE assigned to \mathbf{x} if $\hat{y}(\mathbf{x}) > 0.5$ (majority is ORANGE)

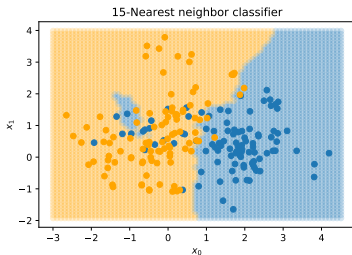
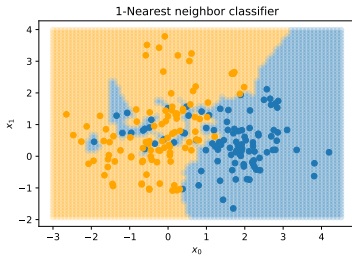
15-Nearest neighbour classifier



Linear classifier vs. 15-Nearest neighbour



1-Nearest neighbour vs. 15-Nearest neighbour



Comparison of techniques

- ▶ At first sight it looks like k-NN has only one parameter, k versus p parameters (number of weights w_i) of the linear model.
- ▶ The **effective** number of parameters of k-NN is N/k which is in general bigger than p (N is the size of the training set).
- ▶ For instance, assume non-overlapping neighbourhoods
 - ▶ There will be N/k neighbourhoods.
 - ▶ To each neighbourhood there correspond one parameter (the mean of the elements of the neighbourhood).

Probability theory

Materials:

- ▶ Chapter 1.3 from **deeplearning**

Probability theory

- ▶ Probability theory is a mathematical framework for dealing with uncertainty, i.e., modeling and analyzing uncertain events and statements
- ▶ In AI probability theory is used in two major ways:
 - ▶ To design AI systems, i.e., derive models and expressions and the corresponding algorithms.
 - ▶ To analyze the behaviour of the AI systems.

Probability theory

- ▶ A **random variable** is a variable that can take values randomly.
- ▶ On its own a random variable just denotes the set of its possible values; to get its full meaning in needs to be coupled with a distribution.

Probability theory

- ▶ There are two types of random variables: **discrete** and **continuous**.
- ▶ Consequently there are two ways to describe probability distributions: **probability mass functions** and **probability density functions**.

Probability mass function

- ▶ The domain of a probability mass function P is the set of all possible states of the random variable x .
- ▶ $\forall x \in \mathcal{X} : 0 \leq P(x) \leq 1$
 - ▶ An impossible event has probability 0 and no state can be less probable than that.
 - ▶ An event that is guaranteed to happen has probability 1 and no state can have a greater chance of occurring.
- ▶ $\sum_{x \in \mathcal{X}} P(x) = 1$
 - ▶ We say that x is **normalized**.
- ▶ Example: Uniform distribution: $P(x = x_i) = \frac{1}{k}$.

Probability density function

- ▶ The domain of the probability density function p must be the set of all possible states of x .
- ▶ $\forall x \in x : p(x) \geq 0$.



$$\int p(x) dx = 1$$

- ▶ Example: uniform distribution $u(x; a, b) = \frac{1}{b-a}$, for $x \in [a, b]$

Conditional probability

- ▶ **Conditional probability** is the probability of some event provided that some other event has happened.
- ▶ Given two random variables x and y , the conditional probability that y has value y provided that we know that x has value x is given by

$$P(y = y \mid x = x) = \frac{P(x,y)}{P(x = x)}$$

- ▶ Another way to see this formula is

$$P(x,y) = P(x = x)P(y = y \mid x = x)$$

i.e., the probability of x and y occurring together is equal to the probability of occurrence of x times the probability of y occurring provided x has occurred.

Expectation

- ▶ The **expectation** or **expected** value of a function $f(x)$ with respect to a probability distribution $P(x)$ is the average value of f over all values x assuming they are drawn from P



$$\mathbb{E}_{x \sim P}[f(x)] = \sum_x P(x)f(x)$$



$$\mathbb{E}_{x \sim P}[f(x)] = \int p(x)f(x)dx$$

- ▶ Linearity of expectations:

$$\mathbb{E}_x[\alpha f(x) + \beta g(x)] = \alpha \mathbb{E}_x[f(x)] + \beta \mathbb{E}_x[g(x)]$$

Variance and covariance

- ▶ The **variance** gives a measure of variation of the values of a random variable x

$$\text{Var}(f(x)) = \mathbb{E}[(f(x) - E[f(x)])^2]$$

Square root of the variance is called **standard deviation**.

- ▶ The **covariance** is a measure of linear relation as well as scale between

$$\text{Cov}(f(x), g(x)) = \mathbb{E}[(f(x) - E[f(x)])(g(x) - E[g(x)])]$$

Covariance matrix

- ▶ The **covariance matrix** of a random vector $\mathbf{x} \in \mathbb{R}^n$ is a $n \times n$ matrix with elements

$$\text{Cov}(\mathbf{x})_{i,j} = \text{Cov}(x_i, x_j)$$

- ▶ The diagonal elements of the matrix give the variance

$$\text{Cov}(x_i, x_i) = \text{Var}(x_i)$$

Bernouli Distribution

- ▶ A distribution over a single binary random variable
- ▶ Controlled by a single parameter $\phi \in [0, 1]$ which corresponds to the probability of the random variable taking the value 1
- ▶ Properties:

$$P(x = 1) = \phi$$

$$P(x = 0) = 1 - \phi$$

$$P(x = x) = \phi^x (1 - \phi)^{1-x}$$

$$\mathbb{E}_x[x] = \phi$$

$$\text{Var}(x) = \phi(1 - \phi)$$

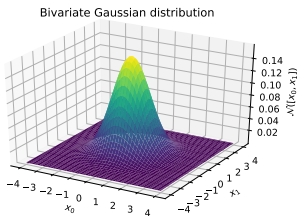
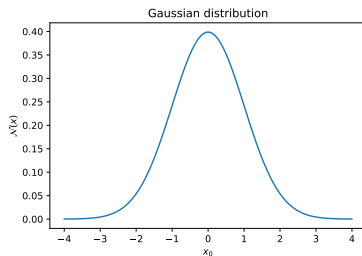
Gaussian distribution

- ▶ The most commonly used distribution, also called **normal distribution**.
- ▶ Controlled by two parameters $\mu \in \mathbb{R}$ (the **mean**) and $\sigma \in (0, \infty)$, (the **standard deviation**)

$$\mathcal{N}(x; \mu, \sigma^2) = \sqrt{\frac{1}{2\pi\sigma^2}} \exp\left(-\frac{1}{2\sigma^2}(x - \mu)^2\right)$$

$$\mathcal{N}(\mathbf{x}; \boldsymbol{\mu}, \boldsymbol{\Sigma}) = \sqrt{\frac{1}{(2\pi)^n \det(\boldsymbol{\Sigma})}} \exp\left(-\frac{1}{2}(\mathbf{x} - \boldsymbol{\mu})^T \boldsymbol{\Sigma}^{-1}(\mathbf{x} - \boldsymbol{\mu})\right)$$

Gaussian distribution

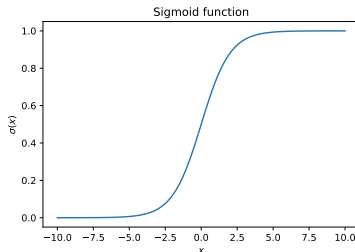


Logistic sigmoid

- ▶ A useful function that we are going to consider

$$\sigma(x) = \frac{1}{1 + \exp(-x)}$$

- ▶ The Logistic (sigmoid) function is commonly used to parametrize Bernoulli distributions.



Bayes' rule

- ▶ Suppose know $P(y | x)$, but we actually need $P(x | y)$. If we know $P(x)$ then we can compute

$$P(x | y) = \frac{P(y | x)P(x)}{P(y)}$$

Although it appears in the formula prior knowledge $P(y)$ is not needed since usually it can be computed as $\sum_x P(y | x)P(x)$

- ▶ It can be straightforwardly derived from the conditional probability formula.
- ▶ It could have be named also after Laplace who independently found it, generalized it, and introduced it in practice.

Acknowledgements

The slides for this lecture were prepared by Mitko Veta and Dragan Bošnjacki.

Some of the slides are based on the accompanying lectures of **deeplearning**.

References